

Probabilistic Exploration for Efficient Bayesian Optimization

Jungwoo Hong

1 Research questions, conclusions, and protocol

This report addresses two research questions. First, does PE help beyond a single objective or acquisition function? Second, can the decay exponent α be related to dimension or selected reliably from a small pilot? Sections 2–4 address the first question, Section 5 addresses the second, and Section 6 tests PE on the Lunar Lander controller objective used by TuRBO.

Main conclusions. PE improves GP-UCB, LogEI, and MES-Gumbel on multiple 30D objectives. With the theory-aligned GP-TS implementation, the benefit repeats on Rosenbrock but not on Bent Cigar or Sparse Hartmann. Useful exponents are usually between 1/2 and 1, yet neither dimension nor a small seed sweep identifies a reliable winner. On Lunar Lander, Decay PE gives a clear and transferable gain with LogEI. In 14D YAHPO tuning, Decay PE helps several UCB, LogEI, and MES settings, mainly by reducing high-loss failures on one dataset, but does not help GP-TS consistently. The current evidence therefore supports PE as a useful option, but does not yet support a universal rule for choosing α .

Protocol. GP signal variance fixed to 1 and not learned, so $k(x, x) \leq 1$; lengthscales learned; GP inputs mapped to $[0, 1]^d$; scrambled Sobol initial design. Within each benchmark, acquisition, and seed all compared policies share the same initial observations and observation-noise stream, so every comparison is paired. The decay family is $p_t = \log(t + 1)/(t + 1)^\alpha$ with t the one-based BO step after the initial design, so a smaller α keeps PE active for more steps. For GP-TS, the previous fixed- β result is replaced throughout by $\beta_t = \log(t + 1)$ and a fresh scrambled-Sobol discretization of size $\text{clip}(\lceil 64\sqrt{t} \rceil, 128, 512)$ at BO step t .

Reading the numbers. Define a higher-is-better performance score as $s = -\text{regret}$ on synthetic benchmarks and $s = \text{return}$ on Lunar. All comparisons then use one convention:

$$\text{improvement over Standard} = s(\text{PE}) - s(\text{Standard}).$$

Equivalently, this is Standard regret minus PE regret on synthetic objectives, and PE return minus Standard return on Lunar. Thus positive always favors PE. Uncertainty is a paired bootstrap 95% confidence interval on that difference. Green favors PE and red favors Standard. A deep, bold cell has an interval that excludes zero; a pale cell does not. Figures use filled markers or outlined heatmap cells for the same distinction. Intervals are not corrected for the simultaneous comparisons in a row, and any cell labelled “best” was selected and evaluated on the same seeds unless a held-out split is stated explicitly.

Experiment	σ_f^2	n_{initial}	BO seeds	Eval. steps
Rastrigin dimension- α sweep, Uniform and MVR	1	$2d$	5	$400 + 2d$
Rastrigin dimension- α sweep, Uniform	1	2	5	402
Rastrigin held-out confirmation of the selected α	1	2	15	402
Rosenbrock and Ackley dimension- α panels	1	2	5	402
30D acquisition sweep, three objectives	1	60	5	500
30D acquisition sweep, two objectives	1	2	5	500
30D acquisition held-out confirmation	1	2	10	500
30D acquisition held-out, exponents 1/2 and 4/3	1	2	10	500
Rastrigin dimension- α sweep at 35 seeds	1	2	35	402
Lunar Lander 12D, four completed acquisitions	1	24	30	400
Lunar held-out controller evaluation	—	—	30	200 episodes/controller
YAHPO XGBoost 14D, three tasks and four acquisitions	1	10	30	120

Table 1: Stages reported here; all passed completeness and protocol audits. Evaluation steps include the initial design. The Lunar held-out row is a separate evaluation of saved controllers, not another BO run.

2 Does PE help across acquisition functions in 30D?

Setup. Three 30D objectives (Sparse Hartmann-6, Rosenbrock, Shifted Bent Cigar) $\times$ four acquisitions (GP-UCB, GP-TS, LogEI, MES-Gumbel) $\times$ six policies: Standard, Decay Uniform at $\alpha = 1/2, 2/3, 1$ and $4/3$, and Fixed $p = 0.2$ as a schedule-free control. LogEI represents the EI family and MES-Gumbel the entropy family. We use 500 evaluation steps and ten paired seeds throughout. Averaged over the completed runs, $\alpha = 1/2$ uses 191 random queries, $\alpha = 2/3$ uses 83, $\alpha = 1$ uses 21, and $\alpha = 4/3$ uses 6, against 97 for Fixed $p = 0.2$.




Figure 1: Endpoint improvement over Standard for every objective, acquisition and PE arm.

Objective	Acquisition	$\alpha = 2/3$ (reference)		$\alpha = 1/2$ (slowest decay)		
		Improv.	Wins	Improv.	95% CI	Wins
Sparse Hartmann-30	GP-UCB	2.140	9/10	2.162	[1.296, 2.941]	10/10
	GP-TS	-0.315	4/10	-0.679	[-1.661, 0.456]	3/10
	LogEI	2.180	9/10	2.373	[1.596, 2.939]	10/10
	MES-Gumbel	2.307	9/10	2.356	[1.535, 2.956]	10/10
Rosenbrock-30	GP-UCB	0.0145	6/10	0.0513	[0.0145, 0.0994]	9/10
	GP-TS	0.0085	7/10	0.0060	[0.0001, 0.0120]	7/10
	LogEI	0.0747	10/10	0.0513	[0.0138, 0.0890]	8/10
	MES-Gumbel	0.0736	9/10	0.0179	[-0.0383, 0.0757]	5/10
Bent Cigar-30	GP-UCB	0.050	4/10	0.474	[0.202, 0.806]	9/10
	GP-TS	-0.0057	5/10	0.0041	[-0.0041, 0.0133]	6/10
	LogEI	0.273	10/10	0.255	[0.163, 0.339]	9/10
	MES-Gumbel	-0.0152	3/10	0.0010	[-0.0191, 0.0216]	5/10

Table 2: The slowest exponent in the sweep against $\alpha = 2/3$.

Findings.

- Three of the four acquisitions have positive mean improvements on all three objectives once the decay is slow enough. LogEI does so at $\alpha = 2/3$; GP-UCB and MES-Gumbel need $\alpha = 1/2$ on at least one objective. Not every individual interval excludes zero.
- **For GP-UCB, the conclusion changes with the exponent.** At $\alpha = 2/3$ GP-UCB gains only on Sparse Hartmann. At $\alpha = 1/2$ it gains on all three: Rosenbrock moves from 0.0145 with a straddling interval to 0.0513 [0.0145, 0.0994] at 9/10 seeds, and Bent Cigar from 0.050 at 4/10 to 0.474 [0.202, 0.806] at 9/10, a tenfold difference in the same cell.
- The improvement tends to decline as α increases. Among the four decay arms, $\alpha = 1/2$ is strongest in seven of the twelve objective-acquisition cells and $\alpha = 4/3$ weakest in eight.
- Under the replacement GP-TS, Rosenbrock improves at $\alpha = 1/2$ (0.0060) and $\alpha = 2/3$ (0.0085), with both paired intervals above zero on the ten held-out seeds. Bent Cigar and Sparse Hartmann have no clear GP-TS gain.
- MES-Gumbel on Bent Cigar is the one block where no arm has a clear gain; every interval contains zero.
- Fixed $p = 0.2$ tracks $\alpha = 2/3$, whose exploration budget it nearly matches: the two share a sign in eleven of twelve cells and their intervals overlap in all twelve. Because their query counts are close but not exactly matched, this comparison alone cannot separate the effect of timing from the total amount of exploration.




Figure 2: Inference regret against evaluation step under LogEI (top) and GP-TS (bottom); lower is better.

Under LogEI the gap opens early on Sparse Hartmann and holds for the rest of the run; on Bent Cigar it opens near step 130 and widens to the end; on Rosenbrock the PE arms are worse than Standard until about step 150. Under GP-TS, Rosenbrock ends below Standard for the two slower decay schedules, whereas Bent Cigar is mixed and Sparse Hartmann does not show a stable gain.

2.1 Replication on unseen seeds

Setup. Sparse Hartmann and Rosenbrock were first run on five pilot seeds, then on ten seeds disjoint from them. $\alpha = 2/3$ is fixed in advance, not selected.




Figure 3: Decay Uniform $\alpha = 2/3$ on the pilot seeds and on the held-out seeds.

Objective	Acquisition	Pilot, seeds 0-4			Held-out, seeds 5-14		
		Improv.	95% CI	W/L	Improv.	95% CI	W/L
Sparse Hartmann-30	GP-UCB	1.838	[0.488, 3.141]	4/1	2.140	[1.266, 2.916]	9/1
	GP-TS	0.607	[-0.035, 1.228]	4/1	-0.315	[-1.182, 0.610]	4/6
	LogEI	1.710	[0.644, 2.776]	4/1	2.180	[1.362, 2.902]	9/1
	MES-Gumbel	1.722	[0.506, 2.937]	4/1	2.307	[1.478, 2.900]	9/1
Rosenbrock-30	GP-UCB	0.1530	[-0.0202, 0.2973]	4/1	0.0145	[-0.0398, 0.0712]	6/4
	GP-TS	0.0158	[-0.0055, 0.0274]	4/1	0.0085	[0.0004, 0.0170]	7/3
	LogEI	0.0681	[0.0391, 0.1027]	5/0	0.0747	[0.0455, 0.1036]	10/0
	MES-Gumbel	0.0013	[-0.0140, 0.0191]	3/2	0.0736	[0.0436, 0.1059]	9/1


Table 3: Pilot and held-out estimates for Decay Uniform $\alpha = 2/3$.

Findings.

- The non-TS cells largely keep their earlier conclusion, and every one significant on the pilot seeds is significant again: Sparse Hartmann GP-UCB $1.838 \rightarrow 2.140$, LogEI $1.710 \rightarrow 2.180$, MES-Gumbel $1.722 \rightarrow 2.307$, Rosenbrock LogEI $0.0681 \rightarrow 0.0747$.
- Rosenbrock MES-Gumbel becomes significant only on the held-out seeds. With the replacement GP-TS, Sparse Hartmann flips from a positive pilot estimate to a negative held-out estimate, with both intervals containing zero. Rosenbrock keeps the positive direction and becomes significant on the held-out seeds (0.0085).
- An earlier pilot with sixty initial evaluations gave a different answer and is not pooled with these results. The conclusions above are for the $n_{\text{initial}} = 2$ protocol shown in Table 1.

3 Does the benefit appear across benchmarks?

Setup. Three functions that extend across dimension — rotated shifted mean-scaled Rastrigin, Rosenbrock, and shifted Ackley — swept over seven exponents at $d \in \{2, 5, 10, 20, 30\}$ under GP-UCB with Decay Uniform. Ackley uses a deterministic interior shift at every dimension so the optimum is never central. Each run has 400 BO steps after two initial evaluations, with five paired seeds.



Cells show PE improvement over Standard for Decay Uniform under GP-UCB, 400 BO steps, 5 paired seeds. Colour is scaled symmetrically about zero, independently per benchmark.

Figure 4: Paired improvement for every dimension and exponent on the three scalable benchmarks.

Benchmark	d	Standard	Best α	PE	Improv.	95% CI	W/L
Rastrigin	2	1.114	1/2	0.301	0.813	[0.503, 1.420]	5/0
	5	1.482	1	1.333	0.149	[0.018, 0.278]	4/1
	10	6.880	2/3	6.192	0.688	[-4.413, 5.789]	2/3
	20	6.367	1/2	2.972	3.396	[-0.801, 7.592]	3/2
	30	9.183	5/6	6.347	2.836	[-0.624, 5.630]	4/1
Rosenbrock	2	0.0027	2/3	0.0050	-0.0023	[-0.0054, 0.0002]	2/3
	5	0.0231	7/6	0.0178	0.0053	[-0.0065, 0.0148]	4/1
	10	0.0895	2/3	0.0596	0.0299	[0.0075, 0.0559]	4/1
	20	0.2815	1/2	0.1289	0.1526	[0.0273, 0.3459]	4/1
	30	0.4435	5/6	0.1864	0.2571	[-0.0154, 0.5948]	4/1
Shifted Ackley	2	0.0531	1/2	0.0175	0.0356	[0.0086, 0.0591]	4/1
	5	1.161	5/6	0.908	0.253	[-0.843, 1.166]	3/2
	10	2.426	1/2	1.870	0.557	[-0.660, 1.718]	3/2
	20	3.027	2/3	2.667	0.360	[-0.253, 1.076]	3/2
	30	5.167	1	3.581	1.586	[0.311, 2.995]	4/1

Table 4: The exponent with the largest mean improvement at each benchmark and dimension.

Findings.

- All three benchmarks contain significant cells in both directions.
- Rosenbrock moves with dimension: three exponents at 2D and $\alpha = 2/3$ at 5D favor Standard, then PE gains at 10D $\alpha = 2/3$ (0.0299 [0.0075, 0.0559]), at 20D for $\alpha = 1/2, 2/3, 5/6$ (0.150 to 0.153), and at 30D $\alpha = 2/3$ (0.2218 [0.0022, 0.5016]).
- Ackley gains at 2D $\alpha = 1/2$ (0.0356 [0.0086, 0.0591]) and at 30D for $\alpha = 2/3, 1, 3/2$ (1.16 to 1.59), but loses at 5D for $\alpha = 2/3$ and $\alpha = 1$.
- Rastrigin gains at 2D $\alpha = 1/2$ (0.813 [0.503, 1.420], 5/5 seeds) and 5D $\alpha = 1$ (0.149), and loses at 20D $\alpha = 1$ (-4.03). Its large 20D and 30D positive means have intervals containing zero.




Figure 5: Inference regret against evaluation step at $d = 5$ and $d = 30$; lower is better.

At $d = 5$ the PE and Standard curves stay inside each other's bands on Rastrigin and Rosenbrock, and on Ackley the $\alpha = 1$ arm stalls above Standard for the whole run. At $d = 30$ the mean PE curves move below Standard inside the first hundred steps and keep the gap to the end on all three benchmarks. The endpoint pattern is therefore not created only by the final few evaluations, although several endpoint intervals still contain zero.

4 Is the best exponent a function of the dimension?

Findings.

- Reading the best exponent down Table 4 gives $1/2, 1, 2/3, 1/2, 5/6$ for Rastrigin, $2/3, 7/6, 2/3, 1/2, 5/6$ for Rosenbrock, and $1/2, 5/6, 1/2, 2/3, 1$ for Ackley. The values cluster in $[1/2, 1]$ and are not ordered by dimension in any benchmark. No dimension has one winner shared by all three benchmarks. Rastrigin and Rosenbrock agree at 10D, 20D, and 30D, while Ackley selects a different exponent in those dimensions.
- On Rosenbrock and Ackley the largest gains appear at higher dimensions, while the best exponent is not monotone in dimension. Every significant high-dimensional cell on those two benchmarks lies in $\alpha \in [1/2, 1]$, except Ackley 30D at $\alpha = 3/2$.
- The largest mean is not always the significant cell: at 30D Rosenbrock the best mean is $\alpha = 5/6$ ($0.2571 [-0.0154, 0.5948]$) while the significant cell is $\alpha = 2/3$.
- Rows of Figure 4 look flat in α wherever PE helps, but that reading comes from five-seed sweeps. On ten held-out seeds (Section 2) the same range is not flat: $\alpha = 1/2$ beats $\alpha = 2/3$ by a factor of ten on Bent Cigar under GP-UCB.
- Figure 5 shows this at the trajectory level: at $d = 30$ the $\alpha = 1/2$ and $\alpha = 1$ curves are close to indistinguishable on all three benchmarks, while $\alpha = 3/2$ sits between them and Standard and keeps part of the gain (+1.48 against +2.57 on Rastrigin, +0.073 against +0.235 on Rosenbrock, +1.16 against +1.59 on Ackley). The flatness holds inside $[1/2, 1]$, not across the whole sweep.

A separate Rastrigin pilot compared seven Uniform and seven maximum-posterior-variance (MVR) decay schedules, plus Standard, using $n_{\text{initial}} = 2d$. Its 375 runs give the same qualitative answer: the preferred exponent is no more predictable from dimension under MVR, and the two point-selection rules can disagree in direction at the same dimension.

5 Can the exponent be chosen in practice?

Setup. The procedure tested is: sweep α on five seeds, keep the per-dimension winner, then evaluate that winner on seeds not used for the selection. The winners from the Rastrigin sweep were rerun against Standard on fifteen disjoint seeds under identical conditions.




Figure 6: Pilot estimate and held-out estimate for the selected exponent at each dimension.

d	Selected α	Pilot improvement	Held-out improvement	95% CI	W/T/L
2	1/2	0.81	0.44	$[-0.31, 1.20]$	12/0/3
5	1	0.15	0.02	$[-0.68, 0.73]$	7/0/8
10	2/3	0.69	-0.53	$[-1.93, 0.88]$	1/5/9
20	1/2	3.40	0.29	$[-2.63, 3.21]$	9/0/6
30	5/6	2.84	0.66	$[-1.69, 3.02]$	9/2/4

Table 5: Pilot versus held-out improvement for the selected exponent.

Findings.

- None of the five selections replicates: every held-out interval contains zero and the effects shrink three- to twelvefold, with 20D falling from 3.40 to 0.29 and 30D from 2.84 to 0.66.
- At 10D the sign reverses, from 0.69 to -0.53 with 1 win, 5 ties, and 9 losses. Win counts still favor PE at 2D, 20D, and 30D.

5.1 Increasing the selection set from five to twenty seeds

Setup. The sweep was extended to 35 seeds. The per-dimension winner is chosen on seeds 0–19 and evaluated on seeds 20–34, which were not used for selection.




Figure 7: Estimate on the twenty selection seeds and on the fifteen untouched seeds, for the exponent chosen at each dimension.

d	Exponent chosen on seeds 0–19				No selection, $\alpha = 1/2$		
	Pick	Train	Test	Test 95% CI	Wins	Test	Wins
2	1/2	0.536	0.869	[0.375, 1.370]	11/15	0.869	11/15
5	1	0.054	−0.158	[−0.920, 0.456]	8/15	− 0.723	4/15
10	4/3	0.487	0.565	[−1.404, 2.520]	5/15	1.084	10/15
20	5/6	1.559	− 2.445	[−4.508, −0.503]	6/15	−1.049	9/15
30	1/2	2.401	−0.307	[−2.192, 1.850]	6/15	−0.307	6/15

Table 6: The selected exponent and a pre-fixed $\alpha = 1/2$, both scored on seeds 20–34.

Findings.

- Only $d = 2$ replicates, at 0.869 [0.386, 1.382]. At $d = 20$ the selected $\alpha = 5/6$ moves from +1.559 on the selection seeds to −2.445 [−4.495, −0.473] on the untouched ones, a reversal that is itself significant.
- Not selecting does better. Fixing $\alpha = 1/2$ scores −0.025 averaged over the five dimensions against −0.295 for the selection rule, though it is significantly worse than Standard at $d = 5$.
- Averaged over 200 random train/test splits inside seeds 0–19, fixing $\alpha = 1/2$ returns $+0.592 \pm 0.179$ and is positive in 100% of splits, while the best selection rule tested — pick the exponent with the highest win rate — returns $+0.466 \pm 0.228$ in 96%. Picking by mean, median, lower confidence bound or a smoothed argmax all do worse. A post-hoc oracle would return +1.040. Thus the grid contains better choices, but the tested selection rules do not identify them reliably.

d	Seeds 0–19	Seeds 20–34	All 35	95% CI	W/L
2	0.536	0.869	0.679	[0.319, 1.062]	28/35
5	−0.294	−0.723	− 0.478	[−0.834, −0.127]	14/35
10	−0.592	1.084	0.126	[−1.040, 1.352]	16/35
20	1.065	−1.049	0.159	[−1.651, 1.962]	21/35
30	2.401	−0.307	1.240	[−0.125, 2.627]	22/35

Table 7: The same cell, $\alpha = 1/2$, estimated on two disjoint seed blocks.

Table 7 explains why. At $d = 30$ the same policy measures +2.401 on seeds 0–19 and −0.307 on seeds 20–34. Rastrigin at 400 BO steps carries enough seed variance that a twenty-seed block estimate moves by more than the effect being measured, so any per-dimension number from this sweep should be read as specific to its seed block.

The simple run-conditioned rule tested here is also insufficient. A 40-step Standard prefix has no consistent association with the later PE outcome: the Spearman correlation between the prefix improvement rate and the eventual gain at $\alpha = 1/2$ is −0.23, 0.11, 0.07, 0.33, −0.08 across the five dimensions, and the correlation with the seed’s own best exponent never exceeds 0.26 in magnitude. The per-seed best exponent is not concentrated either: at $d = 5$ the modal exponent wins on 4 of 20 seeds and at $d = 30$ on 5 of 20, so there is no stable target to predict.

6 Lunar Lander controller tuning

Setup. This is the 12-parameter Lunar Lander controller problem used by TuRBO. One evaluation flies the controller on 50 fixed terrains and uses their average return. We ran 30 seeds. Policies in the same seed share the terrains and the 24 initial Sobol points, so the comparison is paired. Each run has 400 evaluations.

The GP uses a Matérn kernel, fixes signal variance to 1, and learns only the lengthscales. The completed results use GP-UCB, GP-TS, LogEI, and MES-Gumbel with five policies: Standard, decay $\alpha = 1/2, 2/3, 1$, and Fixed $p = 0.2$. The GP-TS row is the completed 150-trial replacement using $\beta_t = \log(t + 1)$ and the increasing Sobol discretization described above; the previous fixed- β result is not used.

A controller is counted as solved when its average return is at least 200. We report return directly because 200 is a success threshold, not the true optimum.




Figure 8: How much each PE policy changes the final return. Positive values favor PE. Right-hand numbers are solved runs out of 30.

Policy	Mean return	Improvement	95% CI	Solved runs	PE queries
Standard	215.8	—	—	21/30	0.0
$\alpha = 1/2$	261.8	46.0	[21.2, 73.4]	27/30	155.3
$\alpha = 2/3$	249.3	33.5	[1.4, 68.0]	25/30	71.2
$\alpha = 1$	261.5	45.6	[23.2, 71.9]	25/30	17.4
fixed $p = 0.2$	232.0	16.2	[−12.8, 44.9]	22/30	73.3

Table 8: LogEI results. Each improvement is compared with Standard using the same 30 seeds. PE queries exclude the 24 initial evaluations.

Findings.

- LogEI is the clearest case. Standard scores 215.8. Decay $\alpha = 1/2$ scores 261.8 (+46.0), and $\alpha = 1$ scores 261.5 (+45.6). Both confidence intervals exclude zero.
- With LogEI, solved runs increase from 21/30 to 27/30 at $\alpha = 1/2$. A crash gives roughly -100 , so a 50-terrain average over 200 means the controller lands reliably.
- At $\alpha = 1/2$, GP-UCB improves by +29.0 and MES-Gumbel by +23.5. Their confidence intervals still include zero.
- Under the replacement GP-TS, $\alpha = 1/2$ changes mean return by +9.1 with a 95% interval of $[-32.8, +50.1]$. The solved count is 4/30 for Standard and 3/30 for PE, so this setting gives no clear evidence of a GP-TS gain on Lunar.

6.1 When PE is used matters

Decay policy	Mean PE queries	Improvement over Fixed	95% CI	W/L
$\alpha = 1/2$	155.3	29.8	[6.6, 54.9]	19/11
$\alpha = 2/3$	71.2	17.3	[−17.2, 50.6]	17/13
$\alpha = 1$	17.4	29.4	[9.1, 52.2]	19/11

Table 9: Decay schedules compared directly with Fixed $p = 0.2$ under LogEI.

Decay $\alpha = 1$ beats Fixed $p = 0.2$ by +29.4 [+9.1, +52.2]. It uses only 17.4 random queries on average, while Fixed uses 73.3. Therefore, the gain is not simply from using more random queries. Using PE early and then returning to LogEI appears better than spreading random queries across the whole run. A matched-count experiment is still needed to confirm this point.

6.2 The improvement transfers to unseen terrains

The results above score each controller on the 50 episode seeds used during its own BO run. To check whether those results are specific to that set, we took all 150 final LogEI controllers (30 BO seeds and five policies) and evaluated every one on the same 200 new episode seeds. None of these seeds appears among the 1,500 seeds used during optimisation. This is a separate test of the final controllers; BO was not run again. In total, the check contains 30,000 new rollouts.




Figure 9: The LogEI improvement on each run’s original 50 episodes and on a shared set of 200 unseen episodes. The two estimates are nearly unchanged.

Policy	Mean return	Gain vs. Standard (95% CI)	Episodes ≥ 200	Bottom 10%	Solved
Standard	209.9	—	73.1%	48.8	20/30
$\alpha = 1/2$	259.5	49.6 [23.7, 76.9]	91.4%	124.7	27/30
$\alpha = 2/3$	249.8	39.8 [10.6, 71.0]	87.5%	114.2	25/30
$\alpha = 1$	255.7	45.8 [22.8, 72.2]	89.0%	113.5	25/30
Fixed $p = 0.2$	227.3	17.4 [−14.3, 48.4]	78.8%	57.4	22/30

Table 10: Performance on 200 unseen episodes. “Episodes ≥ 200 ” is the fraction of individual episodes whose return reaches 200. “Bottom 10%” is the mean of each controller’s 20 lowest returns. A controller is solved when its mean over all 200 episodes reaches 200.

All three decay schedules remain better than Standard: their paired mean-return confidence intervals stay above zero. The largest average result is $\alpha = 1/2$: mean return rises from 209.9 to 259.5, the fraction of episodes reaching 200 rises from 73.1% to 91.4%, and solved controllers rise from 20/30 to 27/30. Its bottom-10% return also rises from 48.8 to 124.7, so the gain is not limited to easy terrains. The original and held-out improvements are +46.0 and +49.6, respectively. Fixed $p = 0.2$ remains uncertain, while decay $\alpha = 1/2$ beats it by +32.2 [+6.9, +59.9] on the new episodes. This is strong evidence against the explanation that PE only found controllers specialised to the 50 episodes used during optimisation.




Figure 10: Mean return of the currently recommended controller as LogEI accumulates evaluations. The dotted line is the solved threshold of 200.

7 Hyperparameter optimisation

As a first transfer test outside analytic functions and simulation control, we ran six scikit-learn model-tuning tasks with GP-UCB and LogEI. The confirmation uses 20 paired seeds per cell, 100 evaluations, and the same fixed unit signal variance. Across 36 acquisition-task-PE comparisons, no PE policy gives a significant held-out test-loss improvement over Standard. The result is useful as a negative transfer check: the current evidence does not support claiming that PE helps every HPO problem.

The higher-dimensional YAHPO study uses GP-UCB, LogEI, theory-aligned GP-TS, and MES-Gumbel with Standard, decay $\alpha = 1/2$, decay $\alpha = 1$, and Fixed $p = 0.2$. Every dimension uses ten initial points and thirty paired confirmation seeds. The completed 14D stage covers three XGBoost tasks and 120 evaluations, for 1,440 audited trials. The 28D and 38D stages use 200 and 250 evaluations and remain in progress; no partial result from either stage is interpreted here.

Because YAHPO reports validation log loss, its improvement is Standard loss minus PE loss. Figure 11 shows the paired mean and bootstrap 95% interval under the same positive-is-better convention used elsewhere.




Figure 11: Endpoint improvement over Standard on the completed 14D YAHPO tasks. Black-edged markers have paired 95% intervals excluding zero.

Dataset	Acquisition	$\alpha = 1/2$	$\alpha = 1$	Fixed $p = 0.2$
credit-g	GP-UCB	+0.0092	-0.0078	+0.0043
	LogEI	-0.0010	-0.0029	+0.0128
	GP-TS	-0.0015	+0.0016	+0.0001
	MES-G	+0.0050	+0.0225	+0.0155
car	GP-UCB	+0.0973	+0.0546	+0.1009
	LogEI	+0.0527	+0.0343	+0.0529
	GP-TS	-0.0167	+0.0036	-0.0027
	MES-G	+0.0800	+0.0198	+0.0961
blood transfusion	GP-UCB	+0.0012	+0.0041	+0.0008
	LogEI	+0.0020	-0.0004	+0.0068
	GP-TS	-0.0084	-0.0002	+0.0005
	MES-G	+0.0023	+0.0067	+0.0073

Table 11: Mean endpoint improvement over Standard in 14D YAHPO. Positive values favor PE. Deep, bold cells have paired 95% intervals excluding zero; complete intervals and win/tie/loss counts are included in the public result CSV.

Across the 24 decay-policy comparisons, 5 improve significantly, 1 is significantly worse, and 18 are inconclusive. The clearest pattern is on **car**: $\alpha = 1/2$ improves LogEI by 0.0527 [0.0100, 0.1045], MES-Gumbel by 0.0800 [0.0224, 0.1455], and GP-UCB by 0.0973 [0.0418, 0.1584]. The median is not uniformly improved; the mean gain comes mainly from reducing the high-loss tail. On **blood transfusion**, $\alpha = 1$ improves MES-Gumbel and GP-UCB by 0.0067 and 0.0041, respectively. No comparison is conclusive on **credit-g**. GP-TS shows no reliable gain and is worse with $\alpha = 1/2$ on **blood transfusion** by 0.0084 [0.0052, 0.0117]. Finally, Fixed $p = 0.2$ also improves UCB, LogEI, and MES-Gumbel on **car**; this stage supports problem-dependent explicit exploration, but does not establish that decay always beats a fixed rate.